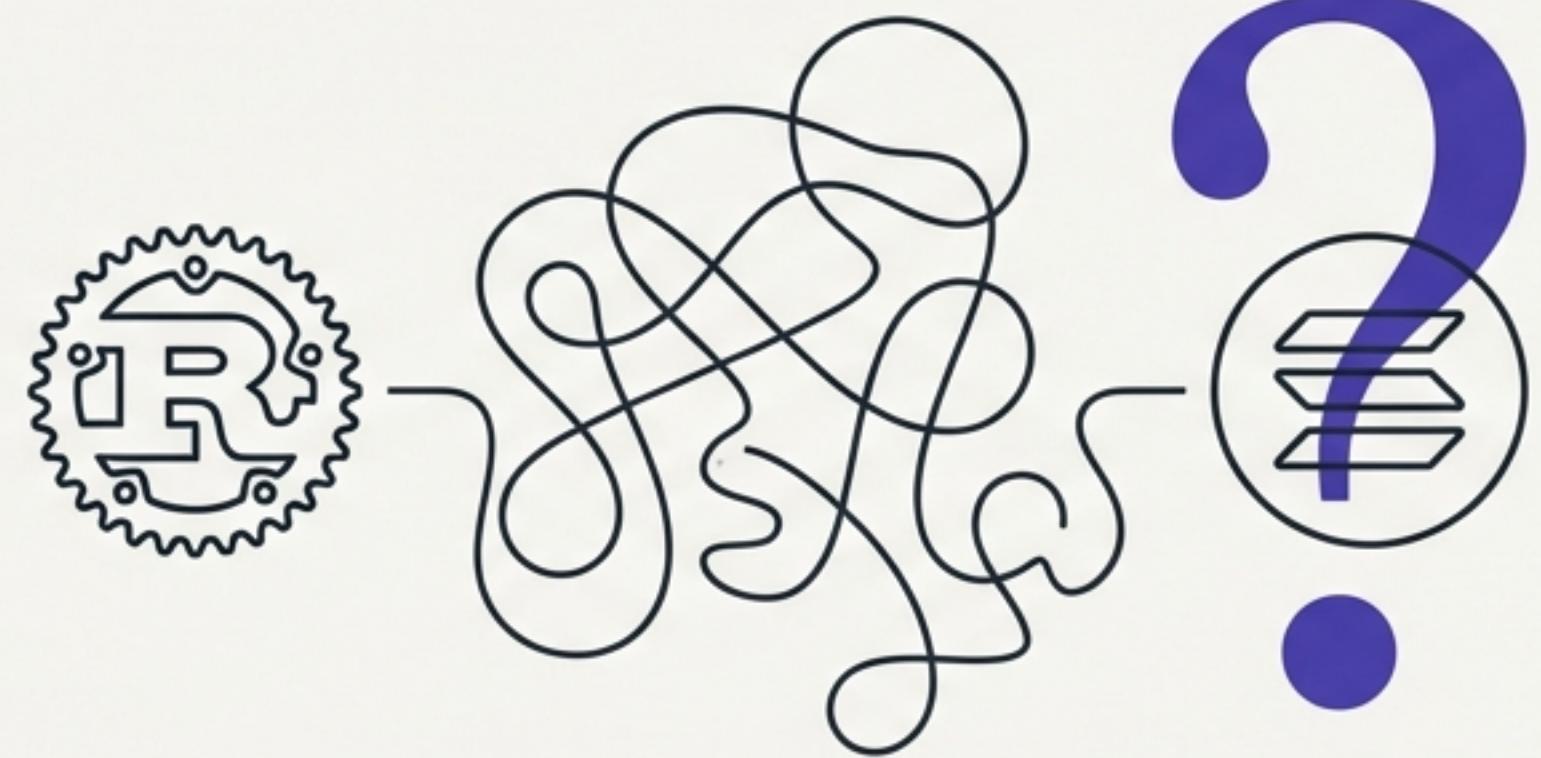


# Why is Native Solana So Confusing for Rust Developers?

You're a proficient Rust developer. You understand ownership, traits, and lifetimes. Yet, when you approach Solana, the rules seem to change.

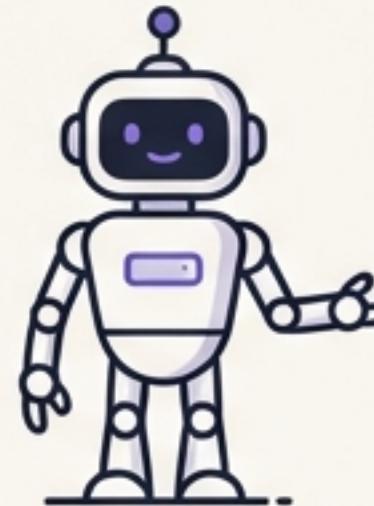
- \* Programs are stateless?
- \* Where does the money actually go?
- \* Why can a program move funds it doesn't 'own'?

This feeling is common. Solana's execution model requires a fundamental mental shift. This deck will provide that shift, moving from a simple analogy to the exact code that powers it.



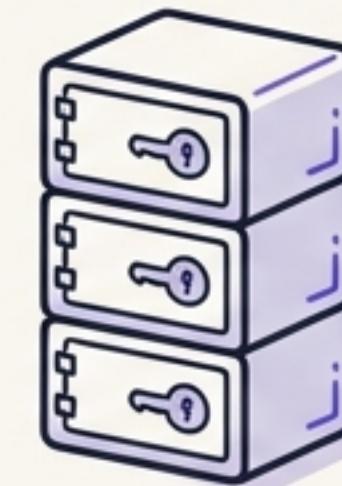
# Let's Think Like a Robot Banker

To understand Solana, forget everything you know about traditional programs for a moment. Instead, imagine a bank run by a very literal robot.



## The Robot Banker

(Your Program) Follows instructions precisely. Has no pockets to hold money.



## The Safe Deposit Boxes

(Accounts) The only place money (SOL) or data can be stored.



## The Ledger

(State) Records who owns what inside the boxes.



## The Cashier

(System Program) The only one who actually moves money between boxes, but only when told.

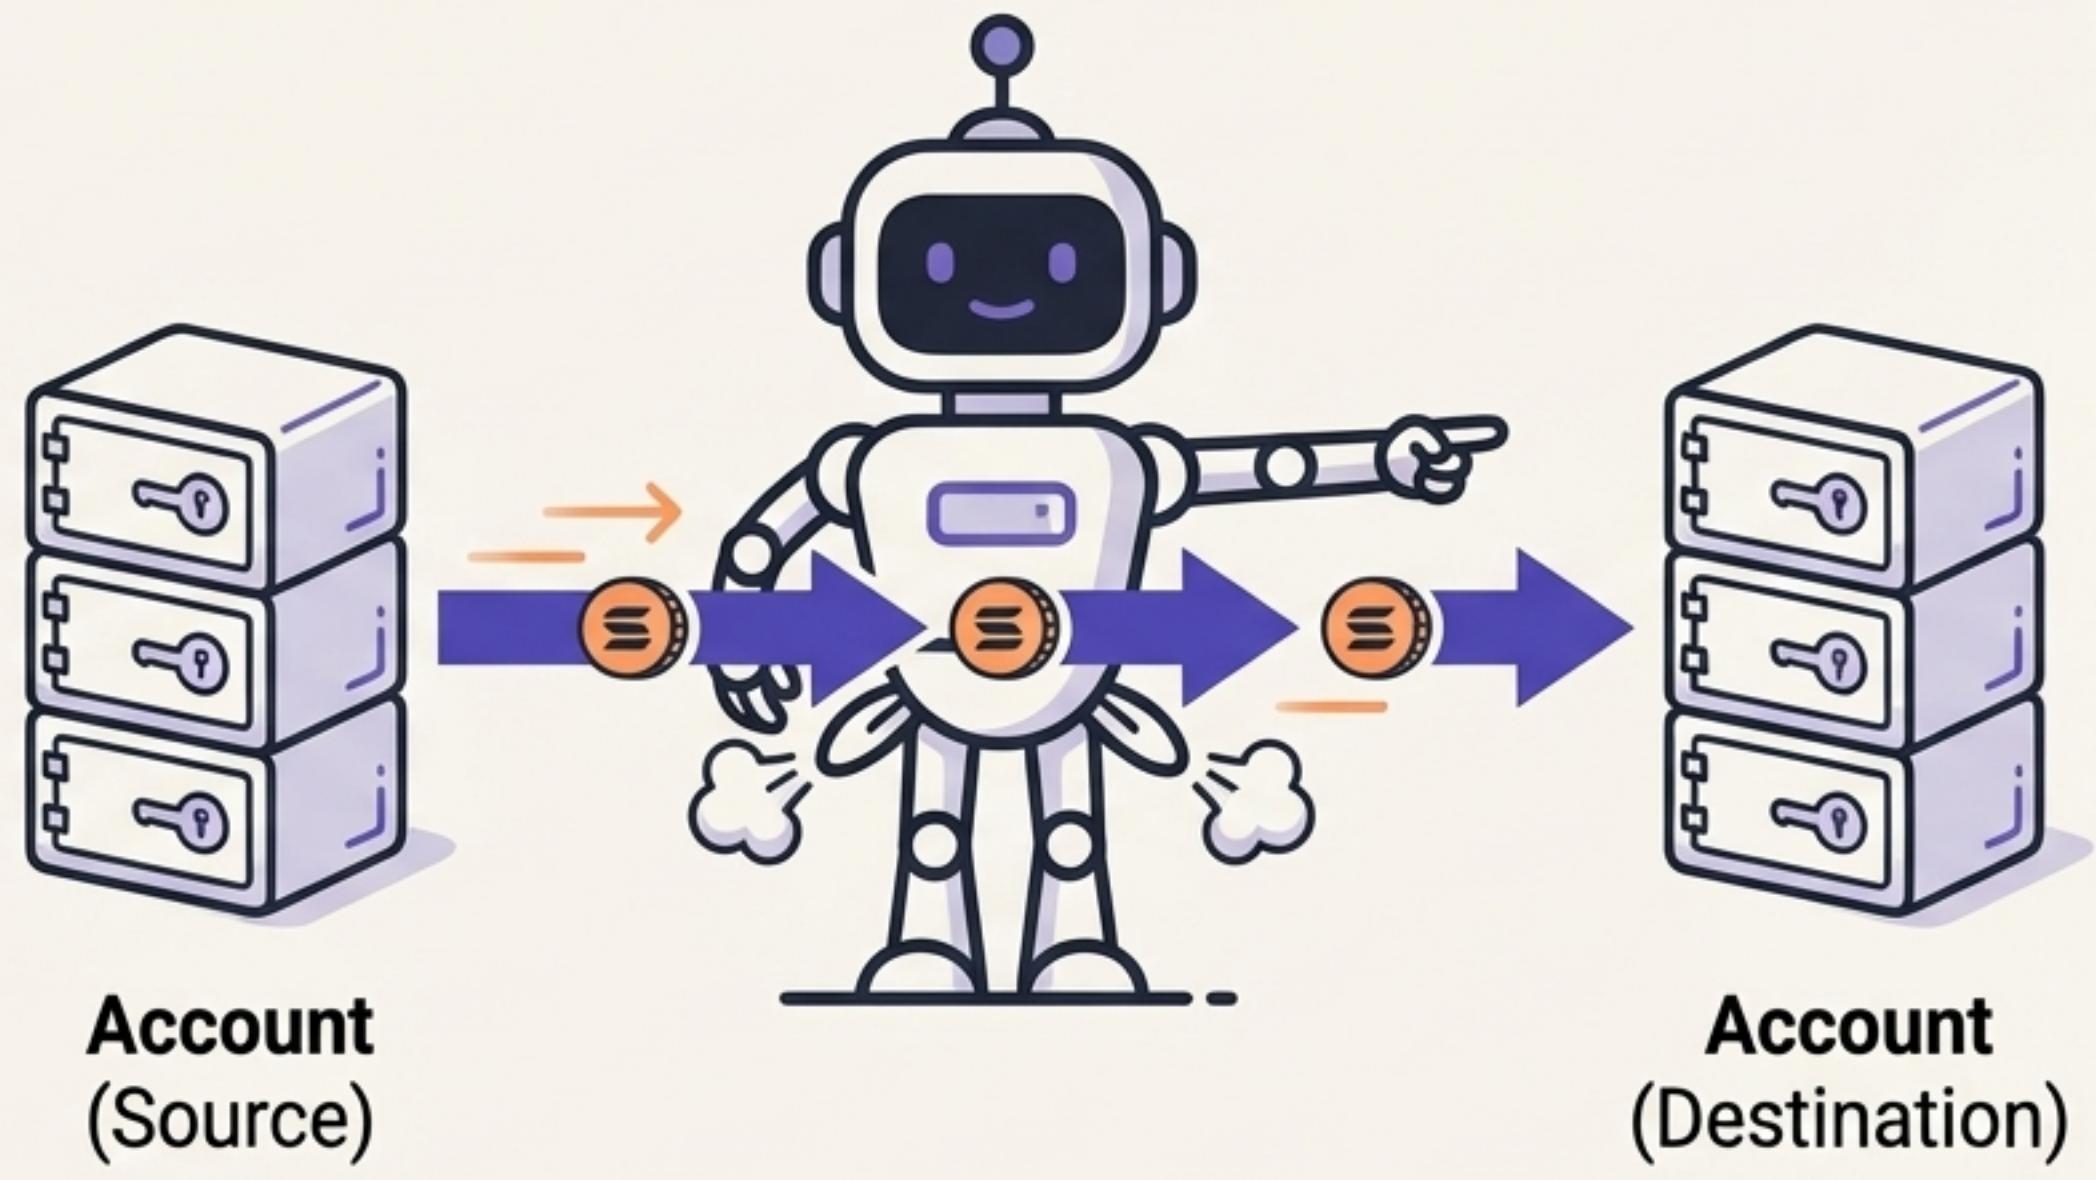
# The Golden Rule: Programs Don't Hold SOL

This is the fundamental truth of Solana that you must internalise.

Your program is the Robot Banker—it is pure logic. It has no wallet, no balance, and no pockets. It can never hold SOL itself.

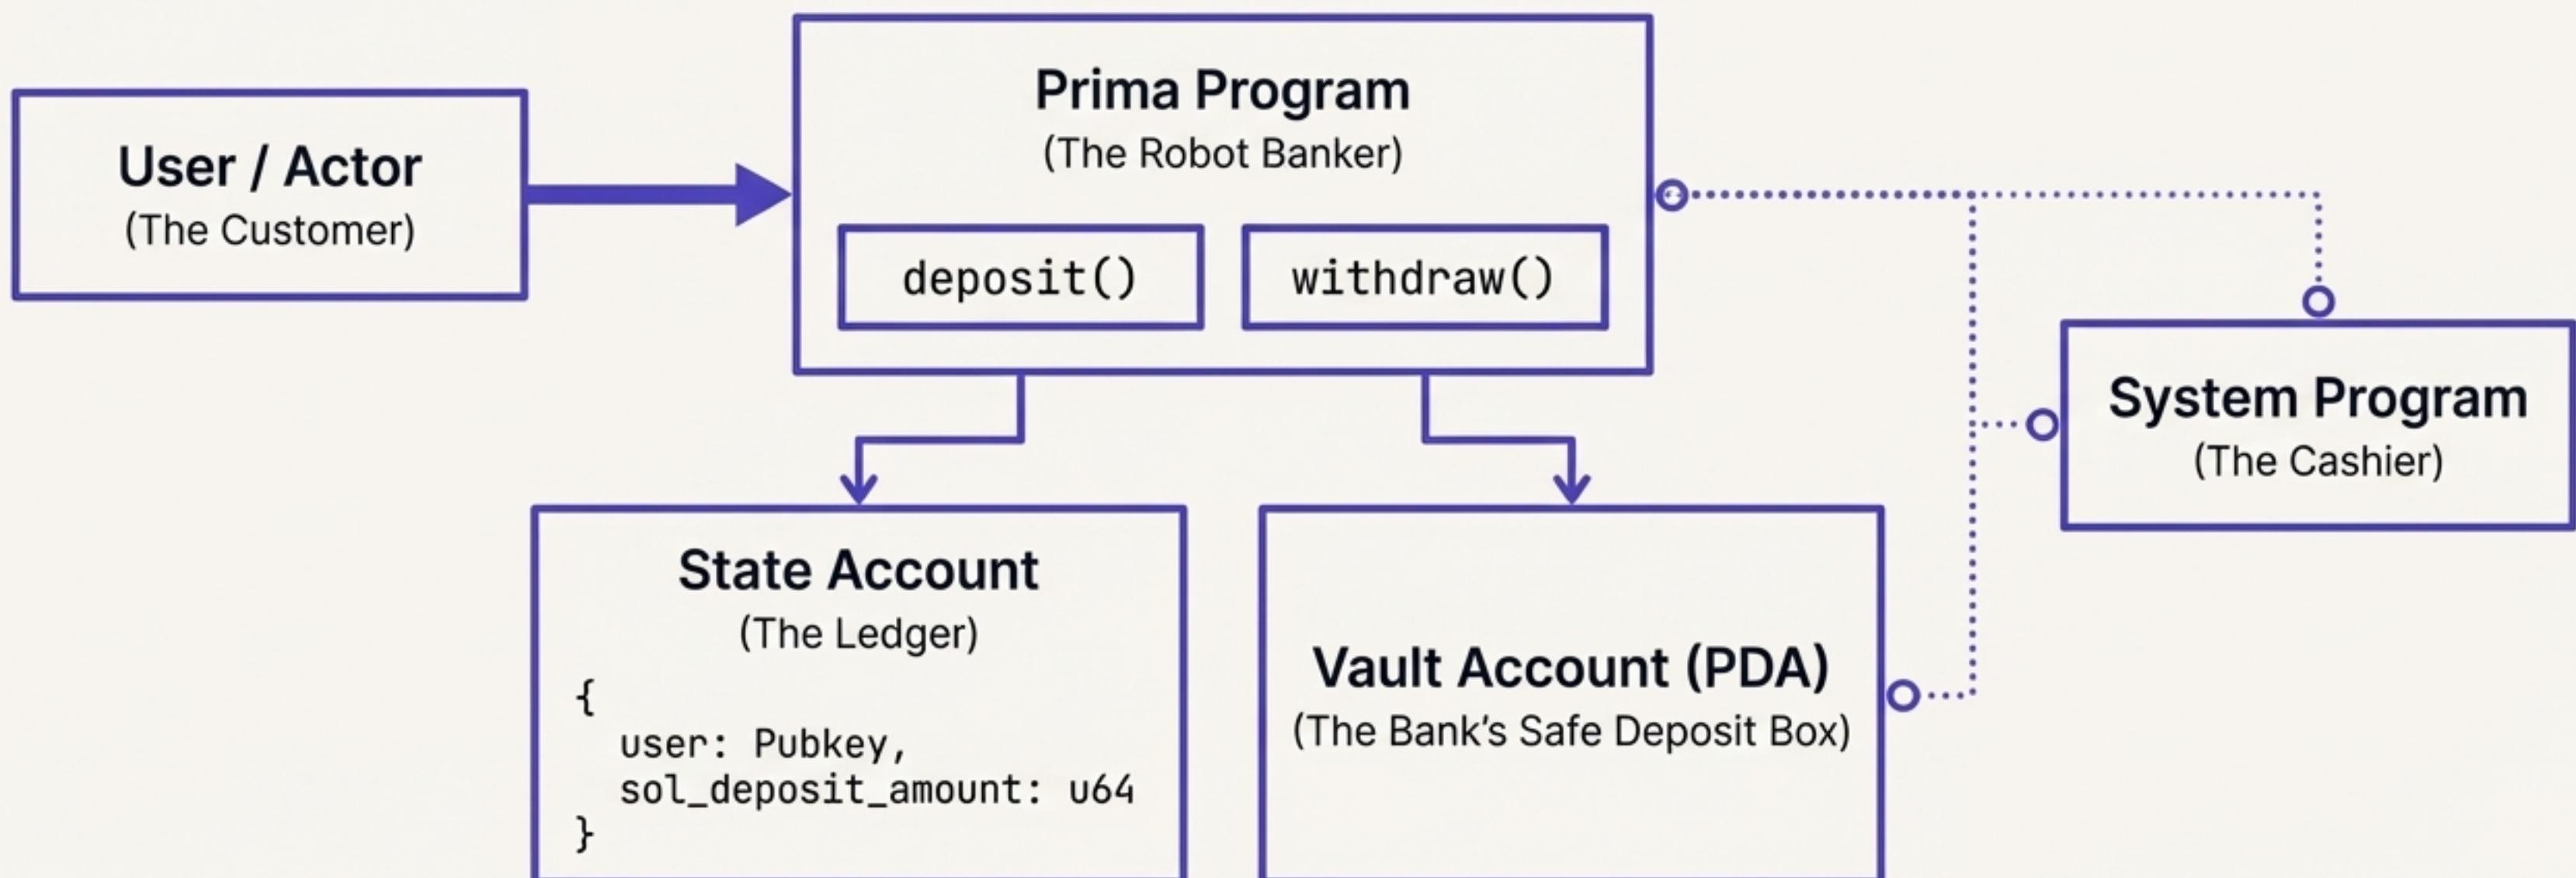
SOL (lamports) only ever lives inside Accounts—the safe deposit boxes.

Your program's primary job is not to *hold* value, but to ***orchestrate the movement of value between accounts.***



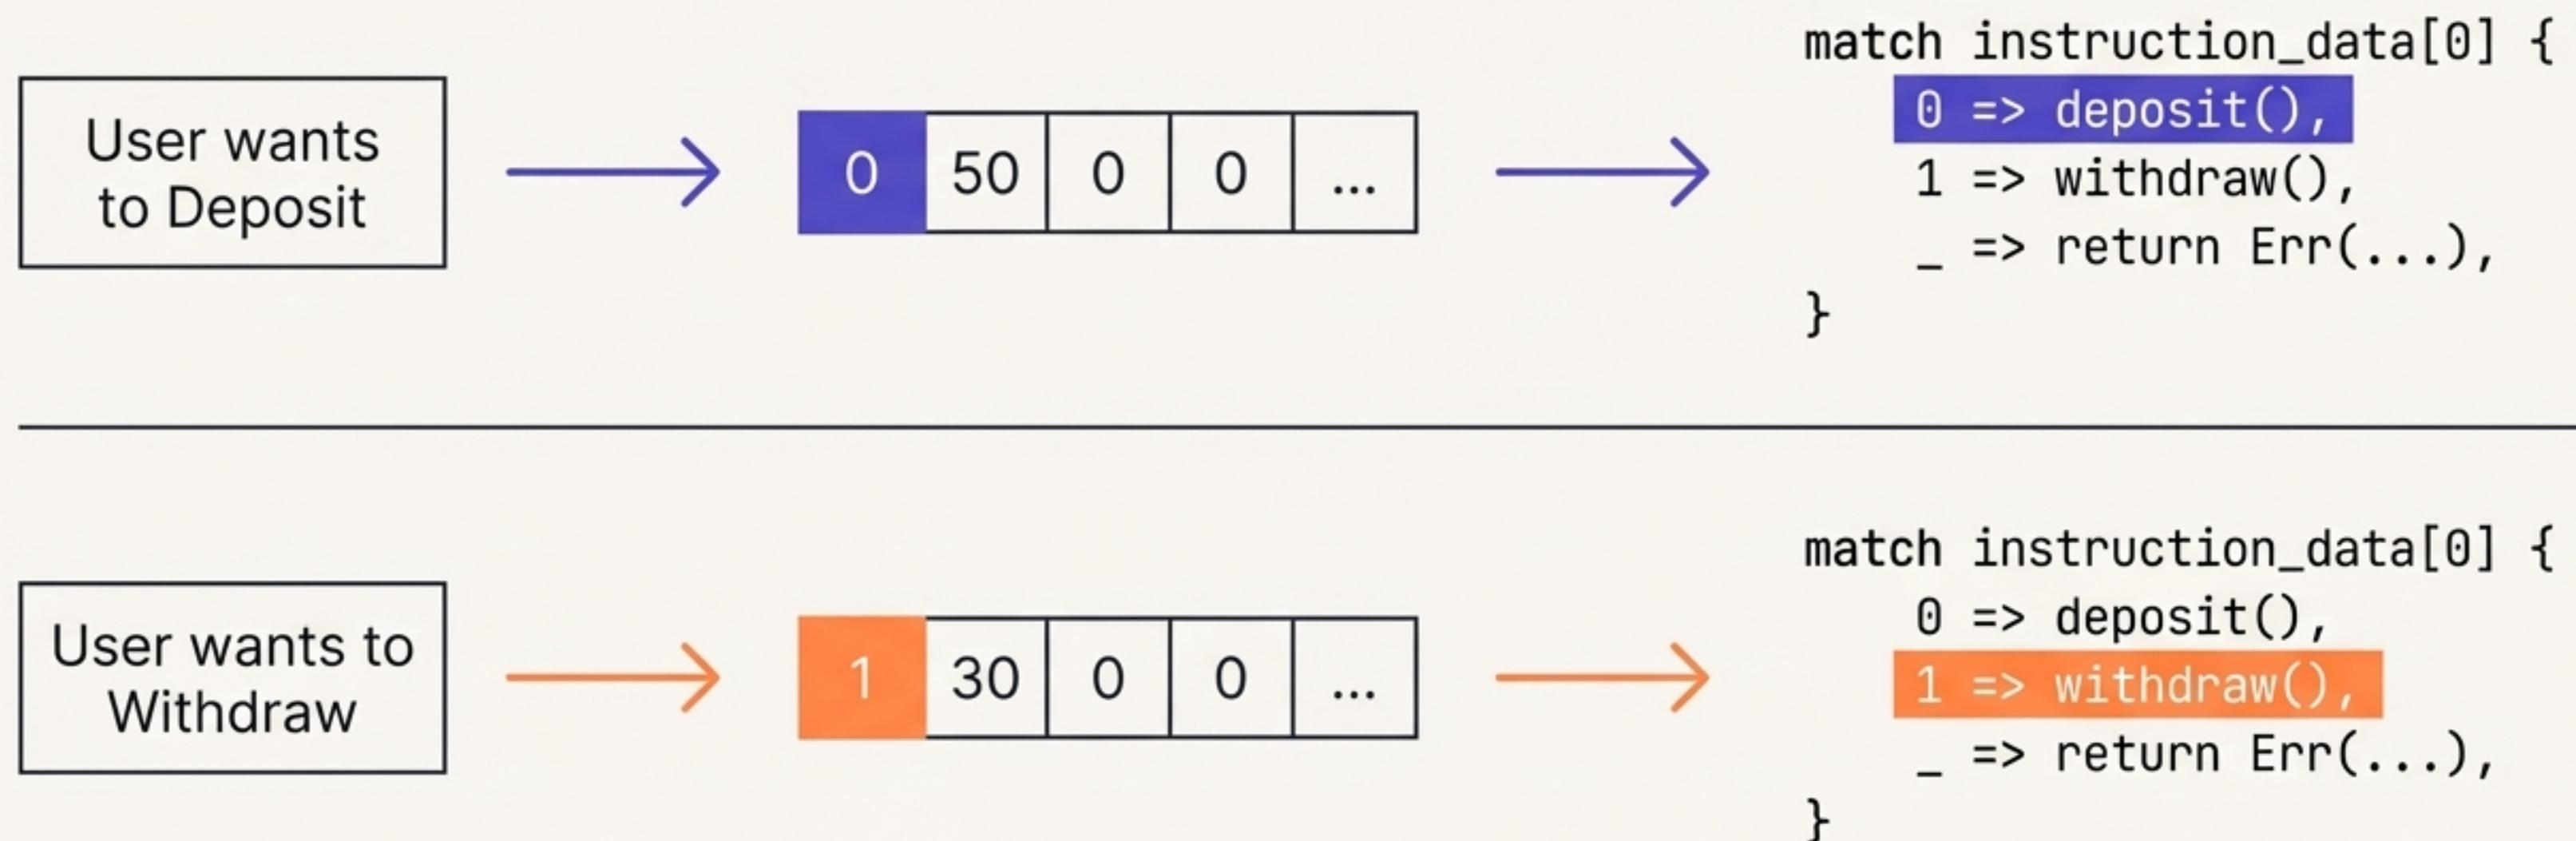
# The Architecture of Our Simple Bank

Our deposit and withdraw contract involves these key players, directly mapping to our analogy. The client (user) must provide all the necessary accounts for the program to work with in every transaction.



# How Does the Robot Know What to Do?

A Solana program is a single entrypoint. It receives raw instruction data as a byte array. We use the very first byte as a ‘discriminant’ to tell the program which function to run. This is our instruction router.



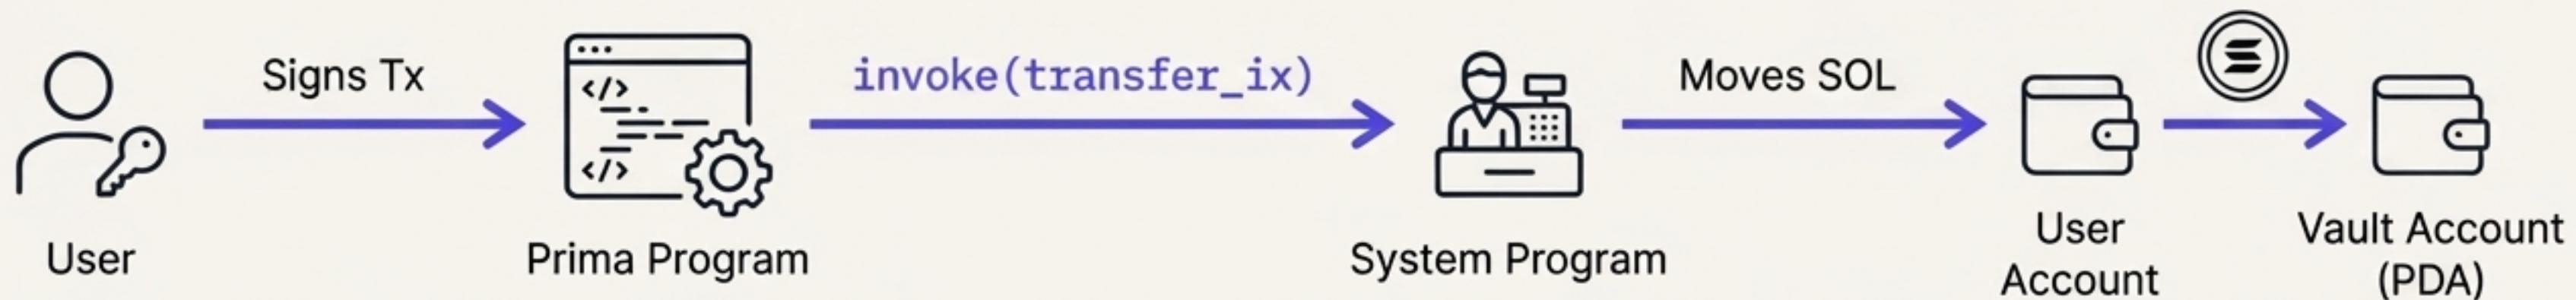
# The Deposit Flow: “Here’s My Money”

When a user deposits SOL, they are the ones authorising the transfer from their own account. They sign the transaction with their private key.

Because a valid signature from the “payer” is already present, our program can simply tell the System Program (“the cashier”) to perform the transfer. This is done using `invoke()`.

Flow:

1. User signs a transaction to call our program's `deposit` instruction.
2. Our program builds a `transfer` instruction (from User to Vault).
3. Our program calls `invoke()`, effectively handing the instruction and the user's signature over to the System Program to execute.



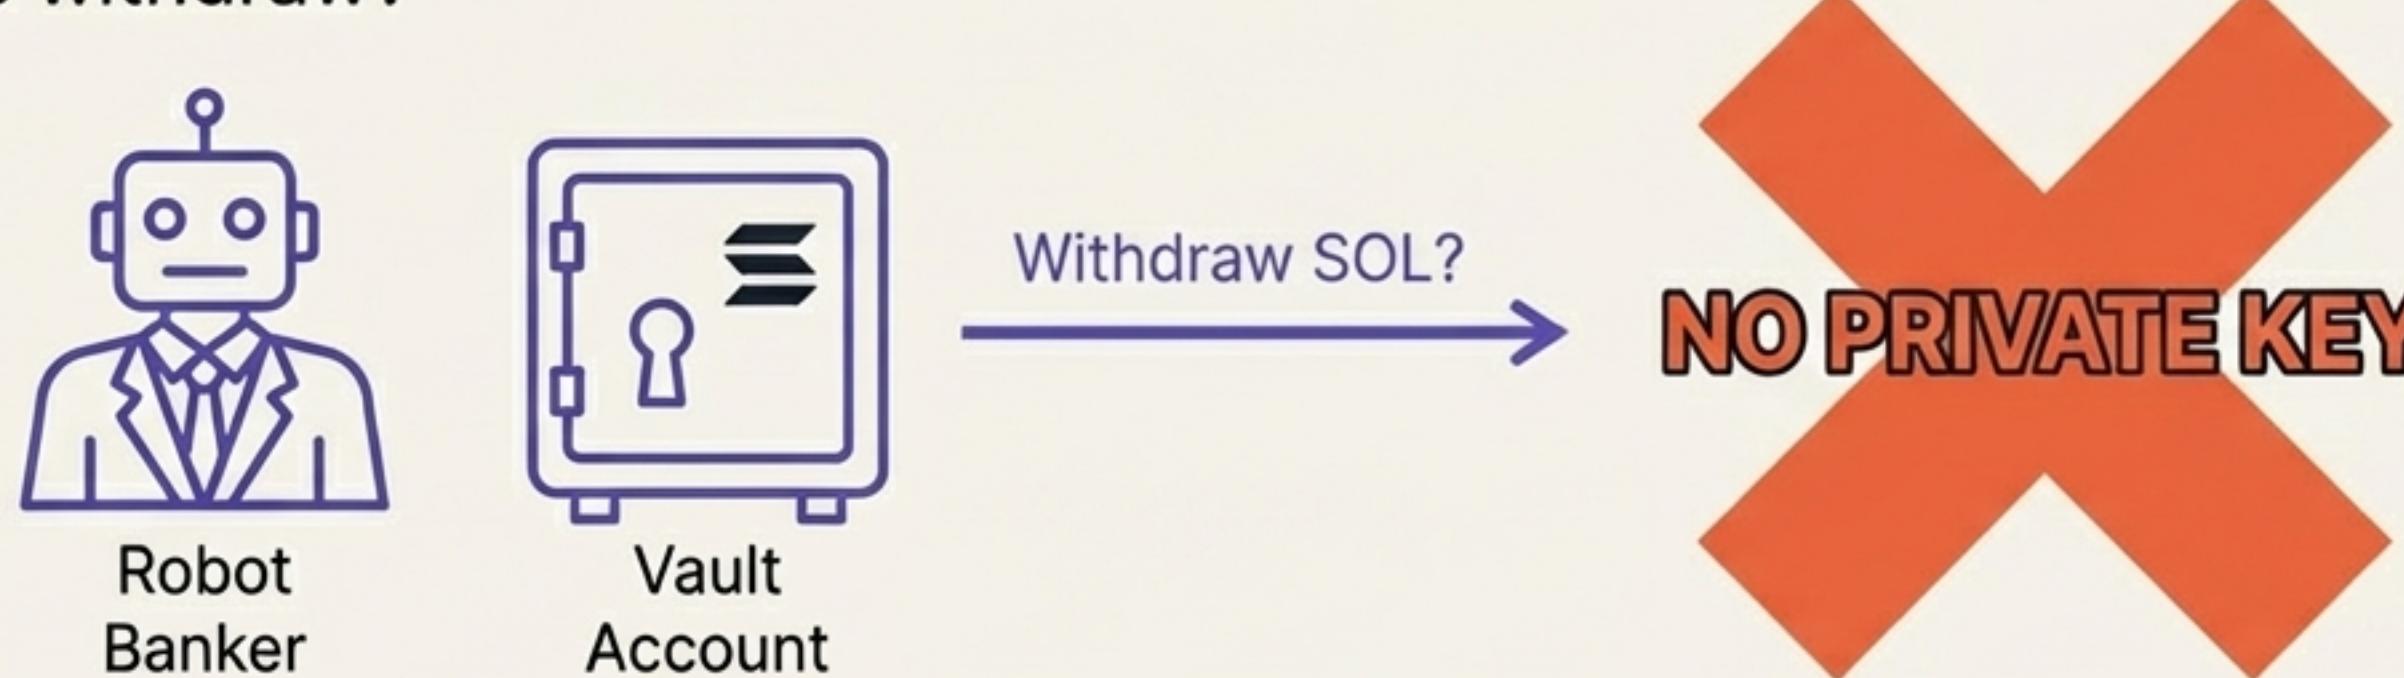
# The Million-SOL Question: Who Signs for the Vault?

Withdrawning is different. The money needs to move *from* the Vault Account to the User.

The problem:

- The Vault Account is controlled by our program.
- But our program is just code—a robot. It has no private key. It cannot ‘sign’ a transaction to authorise the transfer.

If the program can’t sign, how can it ever release the funds? How do we build a bank that lets people withdraw?

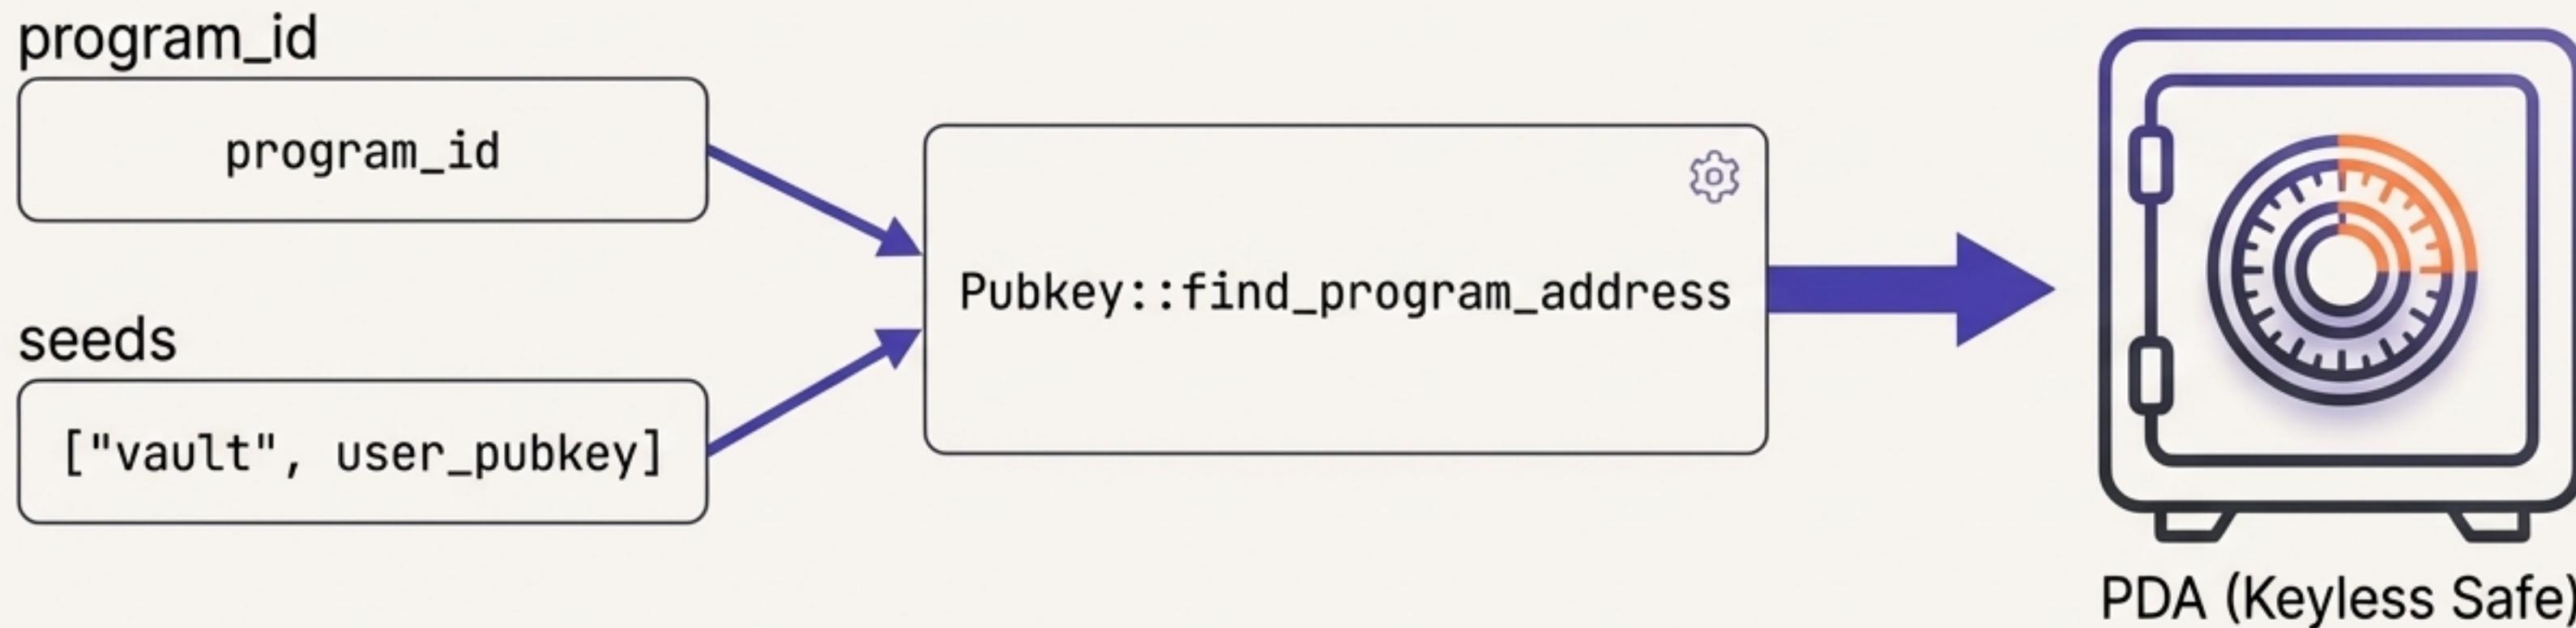


# The Solution: A Keyless Safe Only the Program Can Open

Solana solves this with a **Program Derived Address (PDA)**. A PDA is a special type of account with two key properties:

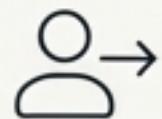
1. It is an address that is **guaranteed to have no corresponding private key**. This means no one can ever directly control it like a normal wallet.
2. It can be 'signed for' *only by the program that created it*.

The program proves its authority by providing the original 'seeds' used to generate the address. Think of the **seeds** as the secret combination to the keyless safe.



# The Two Paths for Moving Value

## Path 1: User Pays (Deposit)

 Who Sends SOL? The User.

 Who Signs? The User, with their wallet's private key.

 How does it work? The user's signature is already on the transaction. The program just needs to "forward" this authority.

 The Function Call: `invoke()`

 Analogy: Handing the cashier a signed cheque.

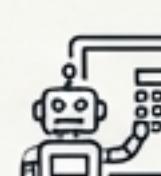
## Path 2: Program Pays (Withdraw)

 Who Sends SOL? The Vault PDA.

 Who Signs? The Program, on behalf of the PDA.

 How does it work? The program provides the PDA's seeds as proof of ownership. Solana runtime verifies these seeds.

 The Function Call: `invoke_signed()`

 Analogy: The robot banker using a secret combination to open the vault for the cashier.

# Inside the Code: The Front Door & The Router

All calls to a native Solana program arrive at a single entrypoint. From there, our `process\_instruction` function acts as a router, inspecting the first byte of the `instruction\_data` to decide which logic to execute.

```
// prima_uccisione/src/lib.rs

use solana_program::...;

// 1. This macro registers our function as the program's single entrypoint.
// It's conditionally compiled to be excluded during unit tests.
#[cfg(not(feature = "no-entrypoint"))]
entrypoint!(process_instruction); ←

// 2. All transactions from the Solana runtime land here.
pub fn process_instruction( ←
    program_id: &Pubkey,
    accounts: &[AccountInfo],
    instruction_data: &[u8],
) -> ProgramResult {
    // 3. We use a match statement on the first byte to route the instruction.
    match instruction_data[0] { ←
        0 => deposit(program_id, accounts, &instruction_data[1..]), // -> Go to deposit logic
        1 => withdraw(program_id, accounts, &instruction_data[1..]), // -> Go to withdraw logic
        _ => Err(ProgramError::InvalidInstructionData),
    }
}
```

1 1. This macro registers our function as the program's single entrypoint. It's conditionally compiled to be excluded during unit tests.

2 2. All transactions from the Solana runtime land here.

3 3. We use a match statement on the first byte to route the instruction.

# Code Deep Dive: The `deposit` Function

The deposit logic is straightforward. It prepares a `transfer` instruction and uses `invoke()` to ask the System Program to execute it. This works because the `user` account, which is the source of the funds, has `is\_signer: true` in the `accounts` array passed by the runtime.

```
1 fn deposit(
2     _program_id: &Pubkey,
3     accounts: &[AccountInfo],
4     instruction_data: &[u8],
5 ) -> ProgramResult {
6     // 1. Get the accounts in the expected order.
7     let acc_iter = &mut accounts.iter();
8     let user = next_account_info(acc_iter)?; // Payer, is_signer = true
9     let vault_pda = next_account_info(acc_iter)?; // Destination
10
11    // 2. Decode the amount from the rest of the instruction data.
12    let amount = u64::from_le_bytes(...);
13
13    // 3. Call invoke() for a Cross-Program Invocation (CPI).
14    // We don't need to provide any extra "signatures" because
15    // the user's signature is already on the parent transaction.
16    invoke(
17        &system_instruction::transfer(user.key, vault_pda.key, amount),
18        &[user.clone(), vault_pda.clone()],
19    )?;
20
21    Ok(())
22 }
```

1. This section parses the accounts from the input array in the expected order, retrieving the user (payer) and the vault PDA (destination). The `is_signer = true` comment is subtly highlighted with a purple tint.
2. This line decodes the transfer amount from the `instruction_data`, using `u64::from_le_bytes`.
3. The `invoke()` function initiates a Cross-Program Invocation (CPI) to the System Program. It specifies the `transfer` instruction with the payer, destination, and amount. Crucially, it relies on the user's signature already present on the parent transaction, indicated by the comment and the `user.clone()` in the arguments.

# Code Deep Dive: The `withdraw` Function & `invoke\_signed`

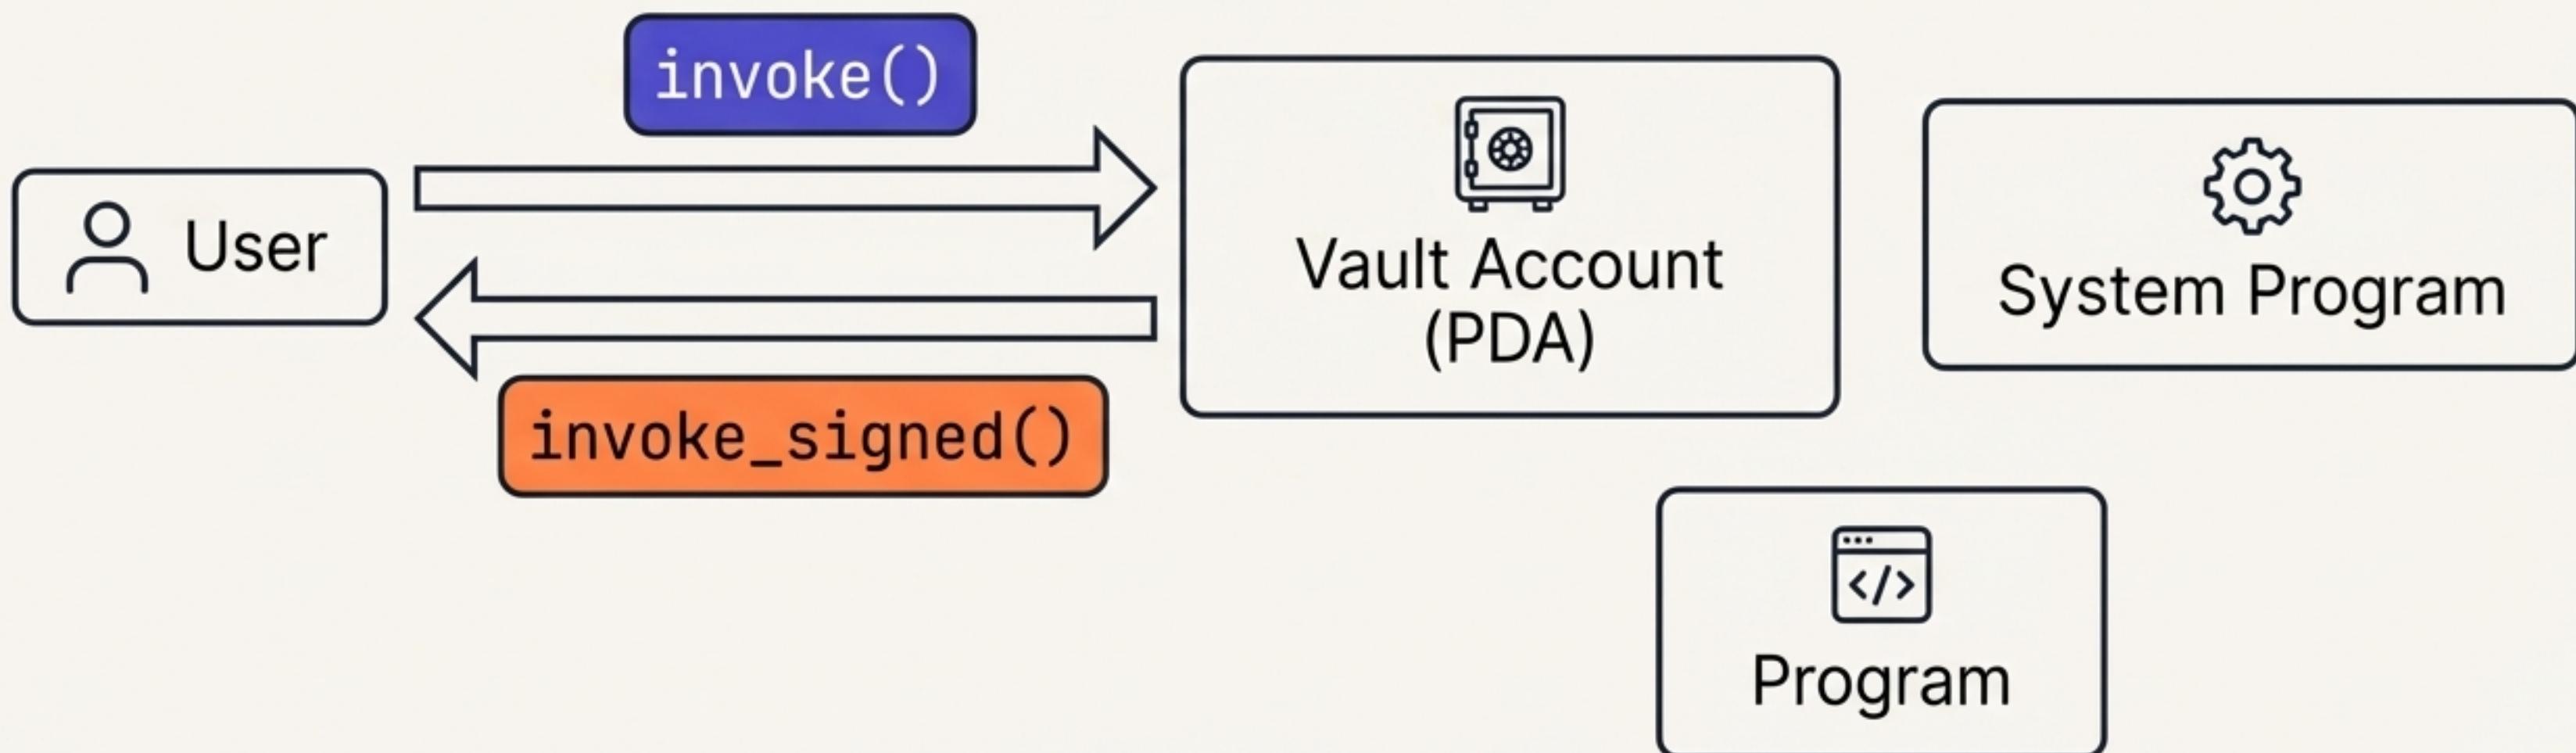
The **withdraw** function is where the magic happens. The program must 'sign' for the Vault PDA to authorise the transfer. It does this by building the PDA's seeds and passing them to `'invoke_signed()'`. The Solana runtime verifies these seeds before allowing the System Program to move the SOL.

```
1 fn withdraw(
2     program_id: &Pubkey,
3     accounts: &[AccountInfo],
4     ...
5 ) -> ProgramResult {
6     // 1. Get accounts. Note the vault_pda is the source of funds.
7     let acc_iter = &mut accounts.iter();
8     let user = next_account_info(acc_iter)?;           // Destination
9     let vault_pda = next_account_info(acc_iter)?;      // Payer, is_signer = false
10
11    // 2. Re-create the PDA seeds. This is the "proof" of ownership.
12    let (expected_vault, bump) = Pubkey::find_program_address(
13        &[b"vault", user.key.as_ref()], // The same recipe
14        program_id,
15    );
16    // Security check: ensure the provided vault account is the correct one.
17    if expected_vault != *vault_pda.key { ... }
18
19    // 3. Build the signer seeds array for the CPI.
20    let seeds = &[b"vault", user.key.as_ref(), &[bump]];
21    let signers = &[&seeds[..]];
22
23    // 4. Call invoke_signed(). This is the program signing!
24    invoke_signed(
25        &system_instruction::transfer(vault_pda.key, user.key, amount),
26        &[vault_pda.clone(), user.clone()],
27        signers, // Pass the seeds as the "signature"
28    );
29    Ok(())
30 }
```

1. **Get accounts.** Parses the 'user' (destination) and 'vault\_pda' (payer, source of funds) from the input array.
2. **Re-create the PDA seeds.** Uses the original seeds ('b"vault"', user's key) and program ID to find the expected PDA and bump. A security check ensures the passed PDA is correct.
3. **Build the signer seeds.** Constructs the array of seeds required for the PDA to 'sign' the CPI, including the crucial bump seed.
4. **Call invoke\_signed().** Executes the cross-program invocation to the System Program to transfer SOL. The 'signers' array provides the PDA's seeds as the program's "signature" for authorization.

# The Complete Picture: Architecture Meets Code

Our architecture diagram now tells the full story. The arrows for moving money are not magic; they are explicit Cross-Program Invocations (CPIs) to the System Program, authorised in two distinct ways.



# What's Really Inside the `accounts` Array?

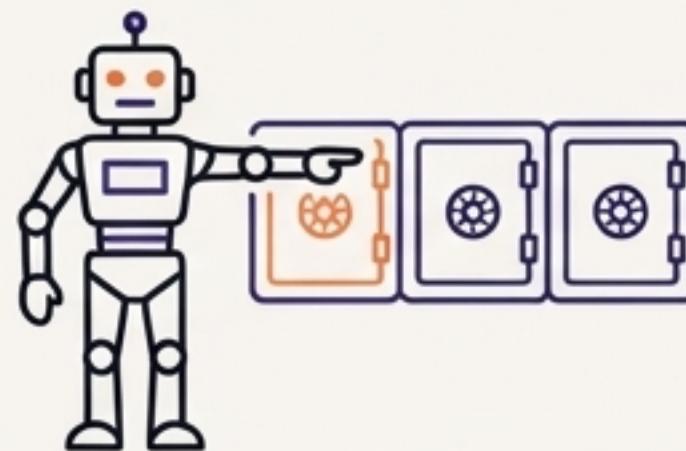
The `accounts: &[AccountInfo]` slice is not just a list of addresses. It's an array of rich objects provided by the client and validated by the Solana runtime. Your program uses this information to get context and check permissions for every operation.

```
accounts: &[
    AccountInfo {           // accounts[0]: The User
        key: <user_pubkey>,
        is_signer: true,      // <-- CRITICAL: The user signed this transaction.
        is_writable: true,    // <-- Balance can be changed.
        owner: SystemProgram,
        ...
    },
    AccountInfo {           // accounts[1]: The Vault PDA
        key: <vault_pda_pubkey>,
        is_signer: false,     // <-- The PDA did not sign.
        is_writable: true,    // <-- Balance can be changed.
        owner: YourProgram,   // <-- Your program controls this account's data.
        ...
    },
    ... // System Program etc.
]
```

For `withdraw`, `user.is\_signer` is still `true` (they initiated the action), but the program uses `invoke\_signed` because `vault\_pda.is\_signer` is `false`.

# Now, You Can Explain Solana's Core Logic

The confusion around native Solana development stems from a few core, counter-intuitive concepts. Now you have the mental model to navigate them.



## The Golden Rule

Programs are stateless robots; they only orchestrate.

**Accounts** hold all state and all SOL.

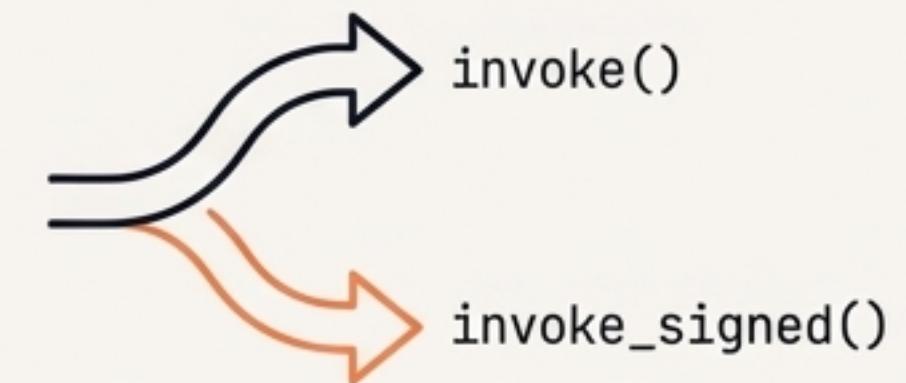


## The Signing Problem

Authority comes from signatures.

**Users** sign with private keys.

**Programs** "sign" for their PDAs using seeds as proof of ownership.



## The Duality of Action

The function you use depends on who is paying.

- `invoke()`: Use when a signer (like a user) is already present in the transaction.
- `invoke_signed()`: Use when a Program Derived Address (PDA) needs to be the signer.